

```

from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

# 3.1 Problem 1. Data Read, Write and Inspect:
# - Dataset for the task: "bank.csv"

# 1. Load the provided dataset and import in pandas DataFrame.
df = pd.read_csv('/content/drive/MyDrive/Workshop/bank .csv')

# 2. Check info of the DataFrame and identify following:

# (a) columns with dtypes=object
object_columns = df.select_dtypes(include=['object']).columns
print(f"Columns with dtype = object:\n {list(object_columns)}\n")

# (b) unique values of those columns.
unique_values = {}
for col in object_columns:
    unique_values[col] = df[col].unique()

print("Unique values for object dtype columns:\n")

for col, values in unique_values.items():
    print(f"{col}: {values}")

# (c) check for the total number of null values in each column.
null_values = df.isnull().sum()
print("\nNull values in each column:\n")
print(null_values)

# 3. Drop columns with dtype=object and store the resulting DataFrame
in "banknumericdata.csv"
df_numeric = df.drop(columns=object_columns)
df_numeric.to_csv('banknumericdata.csv')

# 4. Read "banknumericdata.csv" and find the summary statistics
df_numeric_read = pd.read_csv('banknumericdata.csv')
summary = df_numeric_read.describe()
print("\nSummary statistics:\n")
print(summary)

Columns with dtype = object:
['job', 'marital', 'education', 'default', 'housing', 'loan',
'contact', 'month', 'poutcome', 'y']

Unique values for object dtype columns:

job: ['management' 'technician' 'entrepreneur' 'blue-collar' 'unknown'
'retired' 'admin.' 'services' 'self-employed' 'unemployed']

```

```

'housemaid'
'student']
marital: ['married' 'single' 'divorced']
education: ['tertiary' 'secondary' 'unknown' 'primary']
default: ['no' 'yes']
housing: ['yes' 'no']
loan: ['no' 'yes']
contact: ['unknown' 'cellular' 'telephone']
month: ['may' 'jun' 'jul' 'aug' 'oct' 'nov' 'dec' 'jan' 'feb' 'mar'
'apr' 'sep']
poutcome: ['unknown' 'failure' 'other' 'success']
y: ['no' 'yes']

```

Null values in each column:

```

age          0
job          0
marital      0
education    0
default      0
balance      0
housing      0
loan         0
contact      0
day          0
month        0
duration     0
campaign     0
pdays       0
previous     0
poutcome     0
y            0
dtype: int64

```

Summary statistics:

	Unnamed: 0	age	balance	day
duration \				
count	45211.000000	45211.000000	45211.000000	45211.000000
mean	22605.000000	40.936210	1362.272058	15.806419
std	13051.435847	10.618762	3044.765829	8.322476
min	0.000000	18.000000	-8019.000000	1.000000
25%	11302.500000	33.000000	72.000000	8.000000
50%	22605.000000	39.000000	448.000000	16.000000
100%	45211.000000	45211.000000	45211.000000	45211.000000

75%	33907.500000	48.000000	1428.000000	21.000000
319.000000				
max	45210.000000	95.000000	102127.000000	31.000000
4918.000000				

	campaign	pdays	previous
count	45211.000000	45211.000000	45211.000000
mean	2.763841	40.197828	0.580323
std	3.098021	100.128746	2.303441
min	1.000000	-1.000000	0.000000
25%	1.000000	-1.000000	0.000000
50%	2.000000	-1.000000	0.000000
75%	3.000000	-1.000000	0.000000
max	63.000000	871.000000	275.000000

### # 3.1 Problem 2. Data Imputations

# 1. Load the provided dataset and import in pandas DataFrame

```
df = pd.read_csv('./medical_student.csv')
```

# 2. Check info of the DataFrame and identify columns with missing (null) values

```
print("\nDataFrame Info:\n")
```

```
print(df.info())
```

# Identify columns with missing values

```
missing_values = df.isnull().sum()
```

```
print("\nMissing values in each column:\n")
```

```
print(missing_values[missing_values > 0])
```

# 3. For the column with missing values fill the values using various techniques we discussed above.

```
print("\nFilling missing values:\n")
```

```
for column in missing_values[missing_values > 0].index:
```

```
    if df[column].dtype == 'float64' or df[column].dtype == 'int64':
```

```
        # Fill numeric columns with mean/average value of that column
```

```
        df[column] = df[column].fillna(df[column].mean())
```

```
        print(f"\nFilled missing values in '{column}' with average value of the column.")
```

```
    elif df[column].dtype == 'object':
```

```
        # Fill object columns with mode/most common value from the column
```

```
        df[column] = df[column].fillna(df[column].mode()[0])
```

```
        print(f"\nFilled missing values in '{column}' with most common value in the column.")
```

# 4. Check for any duplicate values

```
print("\nChecking for duplicate rows...")
```

```

duplicates_count = df.duplicated().sum()
print(f"Number of duplicate rows: {duplicates_count}")

# Drop duplicates if any exist
if duplicates_count > 0:
    df = df.drop_duplicates()
    print("Duplicate rows dropped.")

# Final output
print("\nFinal DataFrame Info:")
print(df.info())

```

DataFrame Info:

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200000 entries, 0 to 199999
Data columns (total 13 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   Student ID            180000 non-null float64
 1   Age                   180000 non-null float64
 2   Gender                180000 non-null object
 3   Height                180000 non-null float64
 4   Weight                180000 non-null float64
 5   Blood Type            180000 non-null object
 6   BMI                   180000 non-null float64
 7   Temperature           180000 non-null float64
 8   Heart Rate            180000 non-null float64
 9   Blood Pressure        180000 non-null float64
10   Cholesterol            180000 non-null float64
11   Diabetes               180000 non-null object
12   Smoking               180000 non-null object
dtypes: float64(9), object(4)
memory usage: 19.8+ MB
None

```

Missing values in each column:

Student ID	20000
Age	20000
Gender	20000
Height	20000
Weight	20000
Blood Type	20000
BMI	20000
Temperature	20000
Heart Rate	20000
Blood Pressure	20000

```
Cholesterol      20000
Diabetes          20000
Smoking           20000
dtype: int64
```

Filling missing values:

Filled missing values in 'Student ID' with average value of the column.

Filled missing values in 'Age' with average value of the column.

Filled missing values in 'Gender' with most common value in the column.

Filled missing values in 'Height' with average value of the column.

Filled missing values in 'Weight' with average value of the column.

Filled missing values in 'Blood Type' with most common value in the column.

Filled missing values in 'BMI' with average value of the column.

Filled missing values in 'Temperature' with average value of the column.

Filled missing values in 'Heart Rate' with average value of the column.

Filled missing values in 'Blood Pressure' with average value of the column.

Filled missing values in 'Cholesterol' with average value of the column.

Filled missing values in 'Diabetes' with most common value in the column.

Filled missing values in 'Smoking' with most common value in the column.

Checking for duplicate rows...  
Number of duplicate rows: 12572  
Duplicate rows dropped.

Final DataFrame Info:  
<class 'pandas.core.frame.DataFrame'>  
Index: 187428 entries, 0 to 199999  
Data columns (total 13 columns):

#	Column	Non-Null Count	Dtype
0	Student ID	187428 non-null	float64
1	Age	187428 non-null	float64
2	Gender	187428 non-null	object
3	Height	187428 non-null	float64
4	Weight	187428 non-null	float64
5	Blood Type	187428 non-null	object
6	BMI	187428 non-null	float64
7	Temperature	187428 non-null	float64
8	Heart Rate	187428 non-null	float64
9	Blood Pressure	187428 non-null	float64
10	Cholesterol	187428 non-null	float64
11	Diabetes	187428 non-null	object
12	Smoking	187428 non-null	object

dtypes: float64(9), object(4)

memory usage: 20.0+ MB

None

Cleaned dataset saved as 'cleaned\_medical\_student.csv'.

*# 3.2 Problem 1. Create a DataFrame that is subsetted for the columns 'Name', 'Pclass', 'Sex', 'Age', 'Fare', and 'Survived'.*

*# Retain only those rows where 'Pclass' is equal to 1, representing first-class passengers. What is the mean, median, maximum value, and minimum value of the 'Fare' column?*

```
df = pd.read_csv('Titanic-Dataset.csv')
```

*# Subsetted dataframe*

```
subset_df = df[['Name', 'Pclass', 'Sex', 'Age', 'Fare', 'Survived']]
```

*# filter rows where Pclass is 1.*

```
first_class_df = subset_df[subset_df['Pclass'] == 1]
```

*# Calculate the mean, median, max, and min of the 'Fare' column*

```
mean_fare = first_class_df['Fare'].mean()
```

```
median_fare = first_class_df['Fare'].median()
```

```
max_fare = first_class_df['Fare'].max()
```

```
min_fare = first_class_df['Fare'].min()
```

*# Print the results*

```
print("Problem 1:\n")
```

```
print(f"Mean Fare: {mean_fare}")
```

```
print(f"Median Fare: {median_fare}")
```

```
print(f"Maximum Fare: {max_fare}")
```

```
print(f"Minimum Fare: {min_fare}")
```

*# 3.2 Problem 2. How many null values are contained in the 'Age' column in your subsetted DataFrame?*

*# Once you've found this out, drop them from your DataFrame.*

```

print("\nProblem 2:\n")
# Count the number of null values in the 'Age' column
null_age_count = first_class_df['Age'].isnull().sum()
print(f"Number of null values in 'Age' column: {null_age_count}")

# Drop the rows where 'Age' is null
first_class_df_cleaned = first_class_df.dropna(subset=['Age'])

# Verify that null values are dropped
null_age_count_after_drop =
first_class_df_cleaned['Age'].isnull().sum()
print(f"Number of null values in 'Age' column after dropping:
{null_age_count_after_drop}")

```

Problem 1:

```

Mean Fare: 84.1546875
Median Fare: 60.287499999999994
Maximum Fare: 512.3292
Minimum Fare: 0.0

```

Problem 2:

```

Number of null values in 'Age' column: 30
Number of null values in 'Age' column after dropping: 0

```

*# 3.2 Problem 3. The 'Embarked' column in the Titanic dataset contains categorical data representing the ports of embarkation:*

```

# • 'C' for Cherbourg
# • 'Q' for Queenstown
# • 'S' for Southampton

```

```
df = pd.read_csv('Titanic-Dataset.csv')
```

*# Tasks:*

```

# 1. Use one-hot encoding to convert the 'Embarked' column into
separate binary columns ('Embarked C', 'Embarked Q', 'Embarked S').
embarked_dummies = pd.get_dummies(df['Embarked'], prefix='Embarked')

```

*# 2. Add these new columns to the original DataFrame.*

```
df = pd.concat([df, embarked_dummies], axis=1)
```

*# 3. Drop the original 'Embarked' column.*

```
df.drop(columns=['Embarked'], inplace=True)
```

*# 4. Print the first few rows of the modified DataFrame to verify the changes.*

```
print(df.head())
```

	PassengerId	Survived	Pclass	\
0	1	0	3	

1	2	1	1
2	3	1	3
3	4	1	1
4	5	0	3

SibSp \	Name	Sex	Age
0	Braund, Mr. Owen Harris	male	22.0
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0
1	Heikkinen, Miss. Laina	female	26.0
2	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0
3	Allen, Mr. William Henry	male	35.0

Parch	Ticket	Fare	Cabin	Embarked_C	Embarked_Q
0	A/5 21171	7.2500	NaN	False	False
1	PC 17599	71.2833	C85	True	False
2	STON/O2. 3101282	7.9250	NaN	False	False
3	113803	53.1000	C123	False	False
4	373450	8.0500	NaN	False	False

*# 3.2 Problem 4. Compare the mean survival rates ('Survived') for the different groups in the 'Sex' column.*

```
df = pd.read_csv('Titanic-Dataset.csv')
```

```
mean_survival_by_gender = df.groupby('Sex')['Survived'].mean()
print(f"Mean survival rate by gender:\n{mean_survival_by_gender}")
```

*# Bar plot to show the mean survival rate by gender*


```
plt.figure(figsize=(6, 4))
plt.bar(mean_survival_by_gender.index, mean_survival_by_gender.values,
color=['blue', 'pink'])
plt.title('Mean Survival Rate by Gender')
plt.xlabel('Gender')
plt.ylabel('Mean Survival Rate')
plt.show()
```

Mean survival rate by gender:

```
Sex
female    0.742038
```



```
male      0.188908
Name: Survived, dtype: float64
```



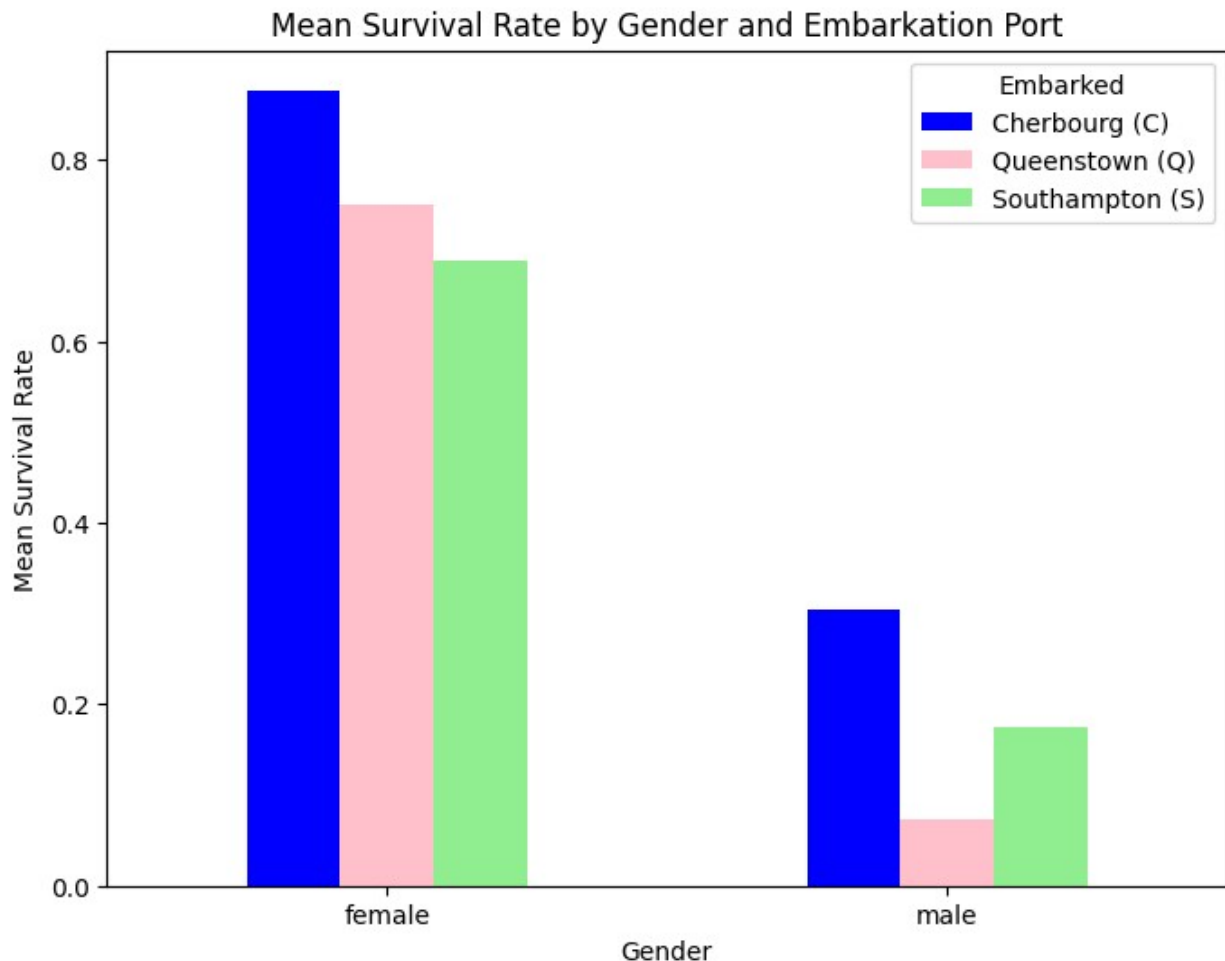
```
# 3.2 Problem 5. Draw a visualization that breaks your visualization
# from Exercise 3 down by the port of embarkation
# ('Embarked'). In this instance, compare the ports 'C' (Cherbourg),
# 'Q' (Queenstown), and 'S' (Southampton).

mean_survival_by_gender_and_embarked = df.groupby(['Sex', 'Embarked'])
['Survived'].mean().unstack()
print(f"Mean survival rate by gender and embarkation port:\n
n{mean_survival_by_gender_and_embarked}")

# Bar plot to show the mean survival rate by gender and embarkation
# port
mean_survival_by_gender_and_embarked.plot(kind='bar', figsize=(8, 6),
color=['blue', 'pink', 'lightgreen'])
plt.title('Mean Survival Rate by Gender and Embarkation Port')
plt.xlabel('Gender')
plt.ylabel('Mean Survival Rate')
plt.xticks(rotation=0) # Rotate x-axis labels to make them readable
plt.legend(title='Embarked', labels=['Cherbourg (C)', 'Queenstown
(Q)', 'Southampton (S)'])
plt.show()
```

Mean survival rate by gender and embarkation port:

Embarked	C	Q	S
Sex			
female	0.876712	0.750000	0.689655
male	0.305263	0.073171	0.174603



# 3.2 Problem 6. (optional) Show how the survival rates ('Survived') vary by age group and passenger class ('Pclass').  
# Break up the 'Age' column into five quantiles in your DataFrame, and then compare the means of 'Survived' by class and age group.  
# Draw a visualization using a any plotting library to represent this graphically.

```
age_quantiles = pd.qcut(df['Age'], 5, labels=['Q1', 'Q2', 'Q3', 'Q4', 'Q5'])  
df['AgeGroup'] = age_quantiles
```

```
mean_survival_by_class_and_age = df.groupby(['Pclass', 'AgeGroup'],  
observed=False)['Survived'].mean().unstack()
```

```

print(f"Mean survival rate by Pclass and Age Group:\n{n{mean_survival_by_class_and_age}}")

# Plot the mean survival rate by Pclass and Age Group
mean_survival_by_class_and_age.plot(kind='bar', figsize=(10, 6),
color=['blue', 'orange', 'green', 'red', 'purple'])
plt.title('Mean Survival Rate by Passenger Class and Age Group')
plt.xlabel('Passenger Class')
plt.ylabel('Mean Survival Rate')
plt.xticks(rotation=0) # Rotate x-axis labels for better readability
plt.legend(title='Age Group', loc='upper left', labels=['Q1', 'Q2',
'Q3', 'Q4', 'Q5'])
plt.show()

```

Mean survival rate by Pclass and Age Group:

AgeGroup	Q1	Q2	Q3	Q4	Q5
Pclass					
1	0.809524	0.761905	0.666667	0.777778	0.506667
2	0.742857	0.400000	0.416667	0.461538	0.363636
3	0.333333	0.197674	0.283582	0.166667	0.088235

